


Introduction to Deep Learning

10. Layers, Blocks, Parameters and GPUs

STAT 157, Spring 2019, UC Berkeley

Alex Smola and Mu Li

courses.d2l.ai/berkeley-stat-157



Caffe

ResNet-101-deploy.prototext

```
layer {  
  bottom: "data"  
  top: "conv1"  
  name: "conv1"  
  type: "Convolution"  
  convolution_param {  
    num_output: 64  
    kernel_size: 7  
    pad: 3  
    stride: 2
```

- Protobuf as the interface
- Good CV model coverage
- Portable
- Not flexible to develop

Tensorflow

Implement Adam

```
# m_t = beta1 * m + (1 - beta1) * g_t
m = self.get_slot(var, "m")
m_scaled_g_values = grad.values * (1 - beta1_t)
m_t = state_ops.assign(m, m * beta1_t,
                        use_locking=self._use_locking)
m_t = state_ops.scatter_add(m_t, grad.indices, m_scaled_g_values,
                            use_locking=self._use_locking)
```

- A domain specific language (DSL) for Python
- A rich set of operators
- Rich features
- Codes are not very easy to read

Keras

```
model = Sequential()
model.add(Dense(512, activation='relu',
               input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))

model.compile(...)
model.fit(...)
```

- Simple DSL for Python, can use multiple backend (TensorFlow, MXNet, CNTK...)
- Easier to use than TensorFlow
- May be slower
- Less convenient to develop and debug

Pytorch

```
class Net(nn.Module):  
    def __init__(self, input_size, hidden_size, num_classes):  
        super(Net, self).__init__()  
        self.fc1 = nn.Linear(input_size, hidden_size)  
        self.relu = nn.ReLU()  
        self.fc2 = nn.Linear(hidden_size, num_classes)  
  
    def forward(self, x):  
        out = self.fc1(x)  
        out = self.relu(out)  
        out = self.fc2(out)  
        return out
```

- Torch tensors + chainer neural networks
- Easy to develop and debug
- Less convenient to deploy

MXNet

Implement Resnet

```
bn1 = sym.BatchNorm(data=data, fix_gamma=False)
act1 = sym.Activation(data=bn1, act_type='relu')
conv1 = sym.Convolution(data=act1, num_filters=64, kernel_size=3, stride=1, pool_size=1, pool_stride=1, num_groups=1)
```

Implement Adam

```
coef2 = 1. - self.beta2**t
lr *= math.sqrt(coef2)/coef1

weight -= lr*mean/(sqrt(variance) + self.epsilon)
```

- Numpy-like Tensor + Keras-like neural networks
- Performance
- Usability

MXNet + Gluon

```
net = gluon.nn.Sequential()
net.add(gluon.nn.Dense(128, activation='relu'))
net.add(gluon.nn.Dense(64, activation='relu'))
net.add(gluon.nn.Dense(10))


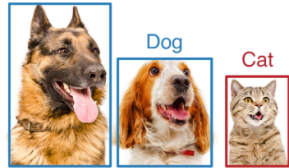
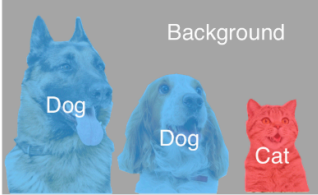
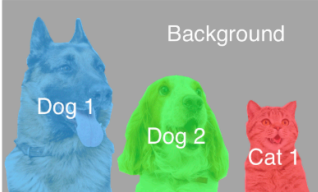
loss = gluon.loss.SoftmaxCrossEntropyLoss()

for data, label in get_batch():
    with autograd.record():
        l = loss(net(data), label)
    l.backward()
    trainer.step(batch_size=data.shape[0])
```

- Numpy-like tensor + Chainer/pytorch-like neural network
- Easy to develop and debug
- Performance

GluonCV

- Toolkit for computer vision
- <https://gluon-cv.mxnet.io/>
- Pre-trained models
- Training scripts to reproduce SOTA results

Application	Illustration	Available Models
<p>Image Classification: recognize an object in an image.</p>	<p>Dog</p> 	<p>50+ models, including ResNet, MobileNet, DenseNet, VGG, ...</p>
<p>Object Detection: detect multiple objects with their bounding boxes in an image.</p>	<p>Dog Dog Cat</p> 	<p>Faster RCNN, SSD, Yolo-v3</p>
<p>Semantic Segmentation: associate each pixel of an image with a categorical label.</p>	<p>Background Dog Dog Cat</p> 	<p>FCN, PSP, DeepLab v3</p>
<p>Instance Segmentation: associate each pixel of an image with an instance label.</p>	<p>Background Dog 1 Dog 2 Cat 1</p> 	<p>Mask RCNN</p>

GluonNLP

- Toolkit for NLP
- <https://gluon-nlp.mxnet.io/>
- Pre-trained models
- Training scripts to reproduce SOTA results

Word Embedding

Mapping words to vectors.

Language Modeling

Learning the distribution and representation of sequences of words.

Machine Translation

From "Hello" to "Bonjour".

Text Classification

Categorize texts and documents.

Sentiment Analysis

Classifying polarity of emotions and opinions.

Parsing

Dependency parsing.

Natural Language Inference

Determine if the premise semantically entails the hypothesis.

Text Generation


Generating language from models.

BERT

Transfer pre-trained language representations to language understanding tasks.

DGL

- Toolkit for graph neural networks
- <https://www.dgl.ai/>
- Relatively new, but with good model coverage
- Has both MXNet and PyTorch backend



Roadmap for 2019

- More toolkits
 - Time series, AutoML, ...
- 100% numpy-compatible
 - A new np package in mxnet
- Compiler integration
 - 50% performance boost on CPU/GPU
 - More hardware coverage: edge, ASIC, ...



Gluon - Imperative Neural Network API

- Create layers and neural networks
 - `gluon.nn`
- Initialize and update parameters
 - `gluon.Parameter`
- Use GPUs